

# ML/DL for Everyone with PYTORCH

## Lecture 7: Wide & Deep

Sung Kim <[hunkim+ml@gmail.com](mailto:hunkim+ml@gmail.com)> HKUST

Code: <https://github.com/hunkim/PyTorchZeroToAll>

Slides: <http://bit.ly/PyTorchZeroAll>

Videos: <http://bit.ly/PyTorchVideo>



# Call for Comments

Please feel free to add comments directly on these slides.

Other slides: <http://bit.ly/PyTorchZeroAll>



# ML/DL for Everyone with PYTORCH

## Lecture 7: Wide & Deep

Sung Kim <[hunkim+ml@gmail.com](mailto:hunkim+ml@gmail.com)> HKUST

Code: <https://github.com/hunkim/PyTorchZeroToAll>

Slides: <http://bit.ly/PyTorchZeroAll>

Videos: <http://bit.ly/PyTorchVideo>



# HKUST PHD Program Application

GPA (a)	Admission?
2.1	0
4.2	1
3.1	0
3.3	1

x\_data = [[2.1],  
[4.2],  
[3.1],  
[3.3]]

y\_data = [[0.0],  
[1.0],  
[0.0],  
[1.0]]



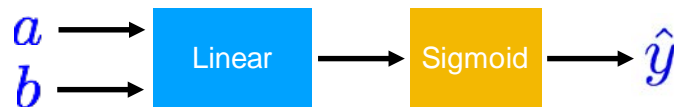
# GPA enough?

## How about experience and others?

GPA (a)	Experience (b)	Admission?
2.1	0.1	0
4.2	0.8	1
3.1	0.9	0
3.3	0.2	1

```
x_data = [[2.1, 0.1],  
           [4.2, 0.8],  
           [3.1, 0.9],  
           [3.3, 0.2]]
```


```
y_data = [[0.0],  
           [1.0],  
           [0.0],  
           [1.0]]
```



# Matrix Multiplication

x\_data = [[2.1, 0.1],  
[4.2, 0.8],  
[3.1, 0.9],  
[3.3, 0.2]]

y\_data = [[0.0],  
[1.0],  
[0.0],  
[1.0]]




$$\underbrace{\begin{bmatrix} a_1 & b_1 \\ a_2 & b_2 \\ \dots & \dots \\ a_n & b_n \end{bmatrix}}_{x \in \mathbb{R}^{N \times 2}} \underbrace{\begin{bmatrix} w_1 \\ w_2 \end{bmatrix}}_{w \in \mathbb{R}^{2 \times 1}} = \underbrace{\begin{bmatrix} y_1 \\ y_2 \\ \dots \\ y_n \end{bmatrix}}_{y \in \mathbb{R}^{N \times 1}}$$

# Matrix Multiplication

x\_data = [[2.1, 0.1],  
[4.2, 0.8],  
[3.1, 0.9],  
[3.3, 0.2]]

y\_data = [[0.0],  
[1.0],  
[0.0],  
[1.0]]



$$\underbrace{\begin{bmatrix} a_1 & b_1 \\ a_2 & b_2 \\ \dots & \dots \\ a_n & b_n \end{bmatrix}}_{x \in \mathbb{R}^{N \times 2}} \underbrace{\begin{bmatrix} w_1 \\ w_2 \end{bmatrix}}_{w \in \mathbb{R}^{2 \times 1}} = \underbrace{\begin{bmatrix} y_1 \\ y_2 \\ \dots \\ y_n \end{bmatrix}}_{y \in \mathbb{R}^{N \times 1}}$$


$$XW = \hat{Y}$$

# Matrix Multiplication

x\_data = [[2.1, 0.1],  
[4.2, 0.8],  
[3.1, 0.9],  
[3.3, 0.2]]

y\_data = [[0.0],  
[1.0],  
[0.0],  
[1.0]]

$$\underbrace{\begin{bmatrix} a_1 & b_1 \\ a_2 & b_2 \\ \dots & \dots \\ a_n & b_n \end{bmatrix}}_{x \in \mathbb{R}^{N \times 2}} \underbrace{\begin{bmatrix} w_1 \\ w_2 \end{bmatrix}}_{w \in \mathbb{R}^{2 \times 1}} = \underbrace{\begin{bmatrix} y_1 \\ y_2 \\ \dots \\ y_n \end{bmatrix}}_{y \in \mathbb{R}^{N \times 1}}$$





$$XW = \hat{Y}$$


```
linear = torch.nn.Linear(2, 1)  
y_prd = linear(x_data)
```




# Go Wide!



# Go Deep!



# Go Deep!

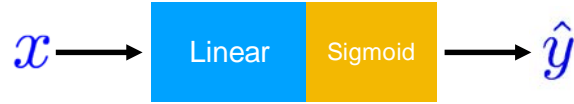


```
sigmoid = torch.nn.Sigmoid()





l1 = torch.nn.Linear(2, 4)
l2 = torch.nn.Linear(4, 3)
l3 = torch.nn.Linear(3, 1)

out1    = sigmoid(l1(x_data))
out2    = sigmoid(l2(out1))
y_pred = sigmoid(l3(out2))
```


# Sigmoid Activation Functions



# Sigmoid: Vanishing Gradient Problem



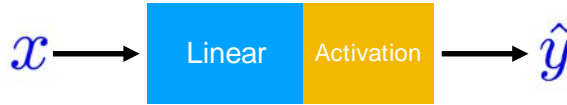
# Activation Functions



Activation function	Equation	Example	1D Graph
Unit step (Heaviside)	$\phi(z) = \begin{cases} 0, & z < 0, \\ 0.5, & z = 0, \\ 1, & z > 0, \end{cases}$	Perceptron variant	
Sign (Signum)	$\phi(z) = \begin{cases} -1, & z < 0, \\ 0, & z = 0, \\ 1, & z > 0, \end{cases}$	Perceptron variant	
Linear	$\phi(z) = z$	Adaline, linear regression	
Piece-wise linear	$\phi(z) = \begin{cases} 1, & z \geq \frac{1}{2}, \\ z + \frac{1}{2}, & -\frac{1}{2} < z < \frac{1}{2}, \\ 0, & z \leq -\frac{1}{2}, \end{cases}$	Support vector machine	
Logistic (sigmoid)	$\phi(z) = \frac{1}{1 + e^{-z}}$	Logistic regression, Multi-layer NN	
Hyperbolic tangent	$\phi(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$	Multi-layer Neural Networks	
Rectifier, ReLU (Rectified Linear Unit)	$\phi(z) = \max(0, z)$	Multi-layer Neural Networks	
Rectifier, softplus	$\phi(z) = \ln(1 + e^z)$	Multi-layer Neural Networks	

Copyright © Sebastian Raschka 2016  
(<http://sebastianraschka.com>)

# Activation Functions



Activation function	Equation	Example	1D Graph
Unit step (Heaviside)	$\phi(z) = \begin{cases} 0, & z < 0, \\ 0.5, & z = 0, \\ 1, & z > 0, \end{cases}$	Perceptron variant	
Sign (Signum)	$\phi(z) = \begin{cases} -1, & z < 0, \\ 0, & z = 0, \\ 1, & z > 0, \end{cases}$	Perceptron variant	
Linear	$\phi(z) = z$	Adaline, linear regression	
Piece-wise linear	$\phi(z) = \begin{cases} 1, & z \geq \frac{1}{2}, \\ z + \frac{1}{2}, & -\frac{1}{2} < z < \frac{1}{2}, \\ 0, & z \leq -\frac{1}{2}, \end{cases}$	Support vector machine	
Logistic (sigmoid)	$\phi(z) = \frac{1}{1 + e^{-z}}$	Logistic regression, Multi-layer NN	
Hyperbolic tangent	$\phi(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$	Multi-layer Neural Networks	
Rectifier, ReLU (Rectified Linear Unit)	$\phi(z) = \max(0, z)$	Multi-layer Neural Networks	
Rectifier, softplus	$\phi(z) = \ln(1 + e^z)$	Multi-layer Neural Networks	

Copyright © Sebastian Raschka 2016  
(<http://sebastianraschka.com>)

[http://rasbt.github.io/mlxtend/user\\_guide/general\\_concepts/activation-functions/](http://rasbt.github.io/mlxtend/user_guide/general_concepts/activation-functions/)

## Non-linear Activations

- ReLU
- ReLU6
- ELU
- SELU
- PReLU
- LeakyReLU
- Threshold
- Hardtanh
- Sigmoid
- Tanh
- LogSigmoid
- Softplus
- Softshrink
- Softsign
- Tanhshrink
- Softmin
- Softmax
- Softmax2d
- LogSoftmax

# Many Activation Functions



David Sheehan

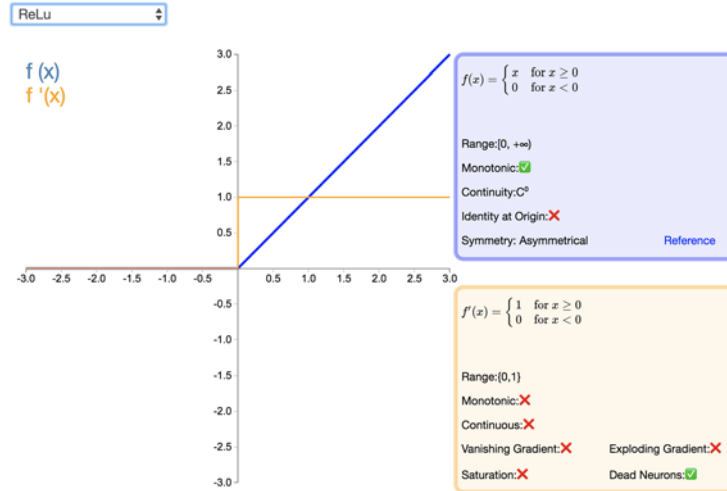
Data scientist interested  
in sports, politics and  
Simpsons references

📍 London via Cork

✉ Email

🐙 Github

Select an activation function from the menu below to plot it and its first derivative. Some properties relevant for neural networks are provided in the boxes on the right.







# Classifying Diabetes



data-diabetes.csv

$X_1$	$X_2$	$X_3$	$X_4$	$X_5$	$X_6$	$X_7$	$X_8$	$y$
-0.411765	0.165829	0.213115	0	0	-0.23696	-0.894962	-0.7	1
-0.647059	-0.21608	-0.180328	-0.353535	-0.791962	-0.0760059	-0.854825	-0.833333	0
0.176471	0.155779	0	0	0	0.052161	-0.952178	-0.733333	1
-0.764706	0.979899	0.147541	-0.0909091	0.283688	-0.0909091	-0.931682	0.0666667	0
-0.0588235	0.256281	0.57377	0	0	0	-0.868488	0.1	0
-0.529412	0.105528	0.508197	0	0	0.120715	-0.903501	-0.7	1
0.176471	0.688442	0.213115	0	0	0.132638	-0.608027	-0.566667	0
0.176471	0.396985	0.311475	0	0	-0.19225	0.163962	0.2	1



<https://www.dropbox.com/sh/zgqixhhcmbquetn/AABOQYrXK8pGM47N9yjqU4aMa?dl=0>

Click


이름	수정됨
diabetes.csv.gz ☆	2020. 6. 18. 오전 8:01
names_test.csv.gz ☆	2020. 6. 18. 오전 8:01
names_train.csv.gz ☆	2020. 6. 18. 오전 8:01
shakespeare.txt.gz ☆	2020. 6. 18. 오전 8:01





# Classifying Diabetes



- Copy the data to your [/My Drive/Colab Notebooks/data](#) folder


 **드라이브**





 **새로 만들기**


☒

 우선순위

 내 드라이브

 공유 드라이브


 공유 문서함

 최근 문서함


내 드라이브 > Colab Notebooks > data ▾

이름 ↑


▾

**diabetes.csv.gz** 


▾

 names\_test.csv.gz 

▾

 names\_train.csv.gz 

▾

 shakespeare.txt.gz 



# Classifying Diabetes



- Mount your GooleDrive to Colab
- Enter the verification code

```
[15] from torch import nn, optim, from_numpy
import numpy as np
from google.colab import drive

drive.mount('/content/gdrive')

xy = np.loadtxt('/content/gdrive/My Drive/Colab Notebooks/data/diabetes.csv.gz', delimiter=',', dtype=np.float32)
x_data = from_numpy(xy[:, 0:-1])
y_data = from_numpy(xy[:, [-1]])
print(f'XW's shape: {x_data.shape} | YW's shape: {y_data.shape}')
```

➡ Drive already mounted at /content/gdrive; to attempt to forcibly remount, call `drive.mount("/content/gdrive", force_remount=True)`.  
X's shape: `torch.Size([759, 8])` | Y's shape: `torch.Size([759, 1])`

# Wide & Deep



```
class Model(nn.Module):
    def __init__(self):
        """
        In the constructor we instantiate two nn.Linear module
        """
        super(Model, self).__init__()
        self.l1 = nn.Linear(8, 6)
        self.l2 = nn.Linear(6, 4)
        self.l3 = nn.Linear(4, 1)
        self.sigmoid = nn.Sigmoid()

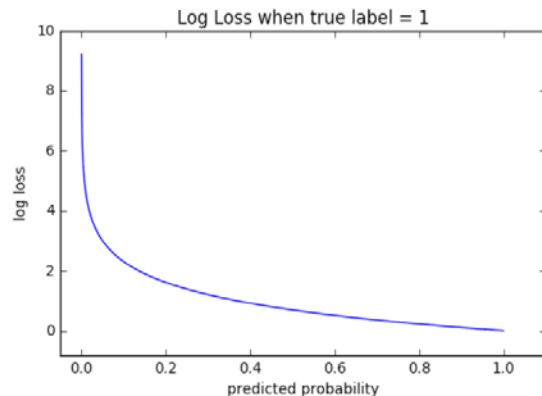
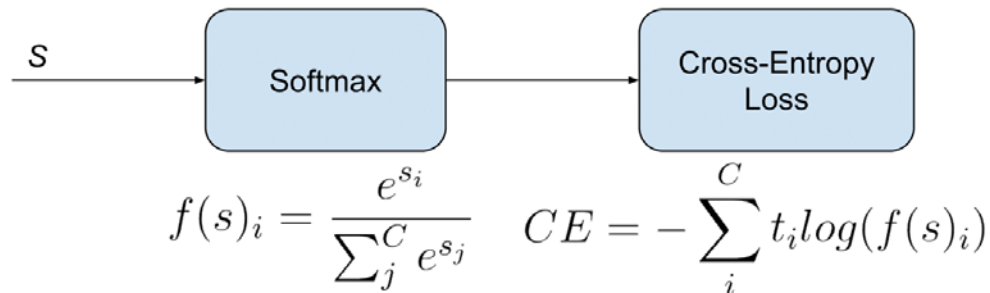
    def forward(self, x):
        """
        In the forward function we accept a Variable of input data and we must return
        a Variable of output data. We can use Modules defined in the constructor as
        well as arbitrary operators on Variables.
        """
        out1 = self.sigmoid(self.l1(x))
        out2 = self.sigmoid(self.l2(out1))
        y_pred = self.sigmoid(self.l3(out2))
        return y_pred
```

```
model = Model()
```

# Wide & Deep



```
# Construct our loss function and an Optimizer. The call to model.parameters()  
# in the SGD constructor will contain the learnable parameters of the two  
# nn.Linear modules which are members of the model.  
criterion = nn.BCELoss(reduction='mean')  
optimizer = optim.SGD(model.parameters(), lr=0.1)
```



# Wide & Deep



```
# Training loop
for epoch in range(100):
    # Forward pass: Compute predicted y by passing x to the model
    y_pred = model(x_data)

    # Compute and print loss
    loss = criterion(y_pred, y_data)
    print(f'Epoch: {epoch + 1}/100 | Loss: {loss.item():.4f}')

    # Zero gradients, perform a backward pass, and update the weights.
    optimizer.zero_grad()
    loss.backward()
    optimizer.step()
```

```
xy = np.loadtxt('data-diabetes.csv', delimiter=',', dtype=np.float32)
x_data = Variable(torch.from_numpy(xy[:, 0:-1]))
y_data = Variable(torch.from_numpy(xy[:, -1]))
```

```
class Model(torch.nn.Module):
```

```
    def __init__(self):
```

```
        """
```

```
        In the constructor we instantiate two nn.Linear module
        """
```

```
        super(Model, self).__init__()
```

```
        self.l1 = torch.nn.Linear(8, 6)
```

```
        self.l2 = torch.nn.Linear(6, 4)
```

```
        self.l3 = torch.nn.Linear(4, 1)
```

```
        self.sigmoid = torch.nn.Sigmoid()
```

```
    def forward(self, x):
```

```
        """
```

```
        In the forward function we accept a Variable of input data and we must return
        a Variable of output data. We can use Modules defined in the constructor as
        well as arbitrary operators on Variables.
        """
```

```
        out1 = self.sigmoid(self.l1(x))
```

```
        out2 = self.sigmoid(self.l2(out1))
```

```
        y_pred = self.sigmoid(self.l3(out2))
```

```
        return y_pred
```

```
# our model
```

```
model = Model()
```

```
# Construct our Loss function and an Optimizer. The call to model.parameters()
```

```
# in the SGD constructor will contain the learnable parameters of the two
```

```
# nn.Linear modules which are members of the model.
```

```
criterion = torch.nn.BCELoss(size_average=True)
```

```
optimizer = torch.optim.SGD(model.parameters(), lr=0.1)
```

```
# Training Loop
```

```
for epoch in range(100):
```

```
    # Forward pass: Compute predicted y by passing x to the model
```

```
    y_pred = model(x_data)
```

```
    # Compute and print Loss
```

```
    loss = criterion(y_pred, y_data)
```

```
    print(epoch, loss.data[0])
```

```
    # Zero gradients, perform a backward pass, and update the weights.
```

```
    optimizer.zero_grad()
```

```
    loss.backward()
```

```
    optimizer.step()
```

# Classifying Diabetes



Design your model using class



Construct loss and optimizer  
(select from PyTorch API)



Training cycle  
(forward, backward, update)

# Output



Epoch: 1/100 | Loss: 0.7010  
Epoch: 2/100 | Loss: 0.6955  
Epoch: 3/100 | Loss: 0.6906  
Epoch: 4/100 | Loss: 0.6861  
Epoch: 5/100 | Loss: 0.6821  
Epoch: 6/100 | Loss: 0.6785  
Epoch: 7/100 | Loss: 0.6752  
Epoch: 8/100 | Loss: 0.6723  
Epoch: 9/100 | Loss: 0.6696  
Epoch: 10/100 | Loss: 0.6672  
Epoch: 11/100 | Loss: 0.6651  
Epoch: 12/100 | Loss: 0.6631  
Epoch: 13/100 | Loss: 0.6613  
Epoch: 14/100 | Loss: 0.6598  
Epoch: 15/100 | Loss: 0.6583

...

Epoch: 85/100 | Loss: 0.6445  
Epoch: 86/100 | Loss: 0.6445  
Epoch: 87/100 | Loss: 0.6445  
Epoch: 88/100 | Loss: 0.6445  
Epoch: 89/100 | Loss: 0.6445  
Epoch: 90/100 | Loss: 0.6445  
Epoch: 91/100 | Loss: 0.6445  
Epoch: 92/100 | Loss: 0.6445  
Epoch: 93/100 | Loss: 0.6445  
Epoch: 94/100 | Loss: 0.6445  
Epoch: 95/100 | Loss: 0.6445  
Epoch: 96/100 | Loss: 0.6445  
Epoch: 97/100 | Loss: 0.6445  
Epoch: 98/100 | Loss: 0.6445  
Epoch: 99/100 | Loss: 0.6445  
Epoch: 100/100 | Loss: 0.6445



# Exercise 7-1: Try to minimize the loss for Diabets data:

- Classify Diabetes with deep nets
  - More than 10 layers
- Find other classification datasets
  - Try with deep network
- Try different activation functions  
Sigmoid to something else
- Find out a better model with respect to minimizing the loss

**WHAT  
NEXT?**



## Lecture 8: DataLoader